

Similar nodes in large graphs

Vlad Sebastian Velici
Supervisor: Adam Prugel-Bennett

October 9, 2014

Given a large graph, in the range of millions or billions of nodes, it is difficult to efficiently find similar nodes. Examples of applications of finding similar nodes in data represented as graphs are all around the web nowadays: *people you may know* in Facebook, *who to follow* or *people similar to you* on Twitter. This project will explore ways to obtain meaningful information of this nature from large datasets represented as graphs.

The goals of this project are developing efficient ways to find similar nodes in large graphs for different use cases and implementing an open-source software library to provide everyone with this functionality.